

Shopware Data Pipeline

Overview

The Shopware Data Pipeline is a robust data engineering solution crafted to collect, process, and analyze data from four sources—two streaming (Web Traffic Logs, CRM Interactions) and two batch (POS Data, Inventory Management Data)—to empower decision-making across Shopware teams. It leverages the Medallion Architecture (Bronze, Silver, Gold layers) to transform raw data into actionable insights, supporting KPI tracking and access via ad-hoc querying, dashboards, and data marts.

High-Level Architecture

Project Objectives

1. **Data Integration:** Seamlessly integrate data from diverse sources, including batch and streaming data.
2. **Data Transformation:** Clean and transform raw data to align with business requirements.
3. **Data Accessibility:** Enable teams to access data through ad-hoc queries, dashboards, or data marts.
4. **Data Storage:** Efficiently organize data using data lakes, warehouses, and marts.
5. **KPI Tracking:** Facilitate tracking of department-specific KPIs.

Medallion Architecture

The pipeline adopts a data lakehouse approach with three layers:

1. **Bronze Layer (Raw Data):** Initial ingestion point for raw data from source systems, chosen for its ability to handle unprocessed data at scale without immediate transformation.
2. **Silver Layer (Processed Data):** Stores cleaned and validated data, selected for its role in providing a reliable intermediate state for further processing.
3. **Gold Layer (Business Insights):** Holds aggregated, analytics-ready data in Redshift and S3, preferred for its optimization for business intelligence and reporting needs.

Data Sources

The pipeline processes data from four sources:

- **POS Data (Batch, Daily):** Sales transactions (quantity, revenue, discounts) for Sales, Operations, and Finance teams, managed in batches for structured, periodic updates.

- **Inventory Management Data (Batch, Hourly):** Real-time inventory and restocking data for Operations and Sales, batched hourly to balance timeliness and processing overhead.
- **Web Traffic Logs (Streaming, Real-Time):** User behavior and session data for Marketing and Data Analysts, streamed for immediate insights into user activity.
- **CRM Interactions (Streaming, Real-Time):** Customer interactions and feedback for Marketing and Customer Support, streamed to enable real-time support and engagement analysis.

Components

Detailed Architecture

Data Ingestion

1. **Batch Data (POS, Inventory):**
 - **S3 Data Lake:** Used as the landing zone for raw batch data in parquet format, chosen for its durability, scalability, and cost-effectiveness for storing large datasets.
 - **Lambda Functions:** Employed to process and move data from Bronze to Silver, selected for its event-driven nature and ability to handle small, specific tasks efficiently.
2. **Streaming Data (Web Traffic, CRM Interactions):**
 - **API Gateway Webhooks:** Implemented to accept pushed data and forward it to Kinesis via Lambda proxies, chosen for its scalability and ease of integrating external data sources.
 - **ECS Fargate Connectors:** Utilized to poll Web Traffic (/api/web-traffic/) and CRM Interactions (/api/customer-interaction/) endpoints, sending data to Kinesis, selected for its containerized, serverless scalability and management simplicity.
 - All streaming data is processed by Lambda functions for real-time computation, preferred for its low-latency execution.

Data Storage & Processing

1. **Kinesis Data Streams:** Manages real-time streaming for Web Traffic and CRM Interactions, chosen for its ability to handle high-throughput data streams with low latency.

2. **Amazon ECS (Fargate):** Runs Docker containers for batch data connectors, selected for its auto-scaling based on CPU/memory usage and seamless integration with Kinesis.
3. **S3 Data Lake:** Stores data across Bronze, Silver, and Gold layers, preferred for its virtually unlimited storage and support for diverse data formats.
4. **AWS Glue:** Handles ETL transformations and metadata management, chosen for its serverless ETL capabilities and integration with the Medallion Architecture.
5. **AWS Lambda:** Processes streaming data and computes KPIs, selected for its cost efficiency and ability to trigger actions in real-time.
6. **EventBridge and Step Functions (Batch Data):** EventBridge triggers Step Functions for Bronze-to-Silver and Silver-to-Gold transformations, chosen for their event-driven automation and orchestration of complex workflows.
7. **Amazon Redshift:** Stores KPI results, preferred for its optimized query performance and integration with Power BI for analytics.

Analytics & Visualization

1. **Amazon Athena:** Enables SQL queries on S3 data, chosen for its serverless query capability without managing infrastructure.
2. **Power BI:** Connects to Redshift for KPI dashboards, selected for its powerful visualization tools and enterprise-grade reporting features.
3. **Data Marts:** Provides team-specific aggregated data, implemented to enhance accessibility and focus on relevant metrics.

Key Performance Indicators (KPIs)

Sales Team (via POS, Inventory Data)

- Total Sales by Region/Product
- Stock Availability
- Product Turnover Rate

Marketing Team (via Web Traffic, CRM Interactions)

- Customer Engagement Score
- Session Duration & Bounce Rate
- Loyalty Activity Rate

Operations Team (via Inventory, POS)

- Inventory Turnover
- Restock Frequency
- Stockout Alerts

Customer Support Team (via CRM Interactions)

- Feedback Score
- Interaction Volume by Type
- Time-to-Resolution

Setup Instructions

Prerequisites

- AWS Account with appropriate permissions
- AWS CLI configured locally
- Docker installed locally
- Terraform installed locally
- Python 3.11 or higher

Deployment Steps

1. **Clone the Repository:**
2. `git clone https://github.com/Amoako419/Shopware.git`

`cd shopware`

3. **Build and Push Docker Images:**
4. `# For CRM Logs Connector`
5. `cd api-gw-webhooks/crm-logs-infra/scripts`
6. `./build_push_ecr.sh`
- 7.
8. `# For Web Logs Connector`
9. `cd api-gw-webhooks/web-logs-infra/scripts`

`./build_push_ecr.sh`

10. **Deploy Infrastructure with Terraform:**
11. `# For CRM Logs Infrastructure`

12. cd api-gw-webhooks/crm-logs-infra/terraform

13. terraform init

14. terraform apply

15.

16. # For Web Logs Infrastructure

17. cd api-gw-webhooks/web-logs-infra/terraform

18. terraform init

terraform apply

19. Deploy Glue Jobs and Step Functions:

- Upload Glue scripts to S3.
- Create and configure Glue jobs and Step Functions using the AWS Console or Terraform.
- Configure EventBridge rules to trigger Step Functions for batch data processing.

20. Set Up Redshift and Power BI:

- Create a Redshift cluster and configure access.
- Connect Power BI to Redshift for dashboard creation.

21. Verify Deployment:

- Check AWS Console to ensure all resources are created.
- Test webhook endpoints for streaming data.
- Monitor CloudWatch logs for connector applications and Step Functions.

Configuration

Key configuration files:

- api-gw-webhooks/crm-logs-infra/terraform/terraform.tfvars: CRM pipeline configuration
- api-gw-webhooks/web-logs-infra/terraform/terraform.tfvars: Web traffic pipeline configuration
- api-gw-webhooks/crm-logs-infra/connector/.env: CRM connector environment variables

- `api-gw-webhooks/web-logs-infra/connector/.env`: Web connector environment variables

Data Flow

Bronze Layer (Data Ingestion)

1. Batch Data Collection (POS, Inventory):

- Lambda moves data from the landing bucket in parquet format:
 - POS: Daily updates.
 - Inventory: Hourly updates.
- Data is sent to Kinesis Data Streams and stored in S3 Bronze buckets.

2. Streaming Data Collection (Web Traffic, CRM Interactions):

- **ECS Fargate Connectors**: Poll web traffic data from `/api/web-traffic/` and CRM interaction data from `/api/customer-interaction/`, sending to Kinesis, chosen for their scalability and reliability.
- **API Gateway Webhooks**: Accept pushed data and forward via Lambda proxies to Kinesis, selected for their flexibility with external integrations.
- Data is processed by Lambda and stored in S3 Bronze buckets.

Silver Layer (Data Processing)

1. Batch Data:

- EventBridge detects new data in the Bronze layer.
- Triggers a Step Function to run AWS Glue jobs for cleaning, validation, and transformation, chosen for their orchestrated workflow management.
- Results are stored in S3 Silver buckets.

2. Streaming Data:

- AWS Lambda processes streaming data in real-time, selected for its low-latency processing.
- Data is cleaned and transformed, then stored in Redshift.

Gold Layer (Analytics)

1. Batch Data:

- EventBridge detects processed data in the Silver layer.

- Triggers a Step Function to compute KPIs using AWS Glue jobs, preferred for its batch processing efficiency.
- Results are stored in Redshift and S3 Gold buckets.

2. **Streaming Data:**

- AWS Lambda computes KPIs for Marketing and Customer Support teams, chosen for its real-time capability.
- Results are stored in Redshift.

3. **Analytics Access:**

- Amazon Athena for SQL queries on S3 data.
- Power BI for dashboards connected to Redshift.
- Notebooks for data science workflows via SageMaker.

Monitoring and Maintenance

1. **CloudWatch Monitoring:** Tracks logs and metrics for Kinesis, Lambda, ECS, and Step Functions, chosen for its comprehensive monitoring and alerting capabilities.
2. **Error Handling:** Implements automatic retries, dead-letter queues, and SNS notifications, selected for their robustness in managing failures.
3. **Scaling:** ECS scales based on usage, Kinesis supports resharding, and Glue adjusts DPU, preferred for their adaptive performance management.

Troubleshooting

Common Issues

1. **Connector Not Sending Data:**

- Check CloudWatch logs for errors
- Verify network connectivity to source APIs
- Ensure IAM permissions are correctly configured

2. **Glue Job Failures:**

- Check job logs in CloudWatch
- Verify input data schema matches expectations
- Check for sufficient IAM permissions

3. **Missing Data in Analytics:**

- Verify Glue crawlers have run successfully
- Check S3 bucket permissions
- Ensure data partitioning is correctly configured

CI/CD with GitHub Actions

Our pipeline uses GitHub Actions for continuous integration and deployment, automating critical processes:

Infrastructure Deployment

- **Terraform Validation:** Validates Terraform configurations on pull requests, chosen to catch errors early.
- **Infrastructure Deployment:** Deploys AWS changes after merging to main, selected for automation efficiency.
- **Security Scanning:** Runs checkov and tfsec, preferred for identifying security vulnerabilities.

Code Quality

- **Python Linting:** Enforces style with flake8 and black, chosen for code consistency.
- **Type Checking:** Validates type hints with mypy, selected for type safety.
- **Unit Tests:** Runs pytest suite, preferred for verifying functionality.
- **Integration Tests:** Tests end-to-end data flow on staging, chosen for system validation.

Container Management

- **Docker Image Building:** Builds images for ECS Fargate connectors, selected for containerized deployment.
- **Image Security Scanning:** Scans with Trivy, chosen for vulnerability detection.
- **ECR Publishing:** Pushes validated images to Amazon ECR, preferred for secure distribution.

Data Quality

- **Schema Validation:** Validates data schemas, chosen to ensure data integrity.
- **Data Test Suite:** Runs dbt tests, selected for transformation accuracy.
- **Documentation:** Auto-generates and publishes data docs, preferred for maintainability.

Workflows are defined in `.github/workflows/` and triggered on:

- Pull request creation/updates
- Merges to main branch
- Scheduled runs for security scanning
- Manual triggers for emergency fixes

Contact

For questions or support, please contact the data engineering team.

Contributing

1. Fork the repository
2. Create your feature branch (`git checkout -b feature/AmazingFeature`)
3. Commit your changes (`git commit -m 'Add some AmazingFeature'`)
4. Push to the branch (`git push origin feature/AmazingFeature`)
5. Open a Pull Request

Code Style

- Python code should follow PEP 8
- Use meaningful variable and function names
- Include docstrings for all functions and classes
- Add type hints where applicable

Testing

- Write unit tests for new features
- Ensure all tests pass before submitting PR
- Include integration tests where necessary

Security

- Never commit sensitive credentials
- Use AWS Secrets Manager for sensitive data
- Follow least privilege principle for IAM roles

Authors

- **Amoako Heskey** - *Contributor*

- **Andrew Marfo** - *Contributor*
- **Charles Adu Nkansah** - *Contributor*
- **Ann-Vanessa Lartey** - *Contributor*

License

This project is licensed under the MIT License - see the LICENSE file for details

Acknowledgments

- AWS Documentation
- Medallion Architecture Best Practices
- Data Engineering Community